

01_delta_basics

August 18, 2026

1 NB1 — Delta Lake Basics (lightweight path)

Stack: `deltalake` (delta-rs) + Polars + DuckDB. No Spark, no JVM. Maps to slide §2 (Delta Lake) + deliverable bullet 1.

Spark equivalent: `spark.read.format("delta").load(path)`
`DeltaTable(path).to_pyarrow_table()`. Same on-disk format, different binding.

```
[1]: import _setup # noqa: F401 -- adds scripts/ to sys.path (file-relative)
import polars as pl
from deltalake import DeltaTable, write_deltalake
from lakehouse import path, reset

table_path = path("scratch", "users_delta")
reset(table_path) # idempotent rerun
```

1.1 1. Write a Delta table

```
[2]: df = pl.DataFrame({
    "id": [1, 2, 3],
    "name": ["alice", "bob", "charlie"],
    "age": [30, 25, 35],
    "city": ["Hanoi", "HCMC", "Danang"],
})
write_deltalake(table_path, df.to_arrow(), mode="overwrite")
```

1.2 2. Read it back + inspect transaction log

Look at `_lakehouse/scratch/users_delta/_delta_log/00000000000000000000.json` — that's the transaction log. Same JSON format Spark/Databricks would write.

```
[3]: dt = DeltaTable(table_path)
print(pl.from_arrow(dt.to_pyarrow_table()))
print("\nHistory:")
for h in dt.history():
    print(f"  v{h['version']} {h['operation']} {h.get('operationMetrics', {})}")
```

shape: (3, 4)

id	name	age	city
---	---	---	---
i64	str	i64	str
1	alice	30	Hanoi
2	bob	25	HCMC
3	charlie	35	Danang

History:

```
v0 WRITE {'execution_time_ms': 3, 'num_added_files': 1, 'num_added_rows': 3, 'num_partitions': 0, 'num_removed_files': 0}
```

1.3 3. Schema enforcement — try to write a wrong schema

```
[4]: bad = pl.DataFrame({"id": [4], "name": ["dan"], "age": ["thirty"], "city": ["Hue"]})
try:
    write_deltalake(table_path, bad.to_arrow(), mode="append")
    print("UNEXPECTED: bad write succeeded - schema enforcement broken")
except Exception as e:
    msg = str(e).splitlines()[0][:120]
    print(f"BLOCKED by schema enforcement (expected): {type(e).__name__}: {msg}")
```

BLOCKED by schema enforcement (expected): Exception: Cast error: Cannot cast string 'thirty' to value of Int64 type

1.4 4. Schema evolution (opt-in)

```
[5]: new = pl.DataFrame({
    "id": [4], "name": ["dan"], "age": [28], "city": ["Hue"], "tier": ["premium"],
})
write_deltalake(table_path, new.to_arrow(), mode="append", schema_mode="merge")
dt = DeltaTable(table_path)
# Sort by id so the printout is stable across reruns - Delta does not
# preserve write-order across appends.
print(pl.from_arrow(dt.to_pyarrow_table()).sort("id"))
```

shape: (4, 5)

id	name	age	city	tier
---	---	---	---	---
i64	str	i64	str	str

1	alice	30	Hanoi	null
2	bob	25	HCMC	null
3	charlie	35	Danang	null
4	dan	28	Hue	premium

1.5 5. Bonus — query with DuckDB (zero copy)

```
[6]: import duckdb

# We hand DuckDB an Arrow table rather than calling `delta_scan()`. delta_scan
# autoloads a DuckDB extension over the network - fine at home, a support
# ticket in a firewalled classroom. Arrow registration is zero-copy and offline.
con = duckdb.connect()
con.register("users", DeltaTable(table_path).to_pyarrow_table())
tier_counts = con.sql("SELECT tier, count(*) AS n FROM users GROUP BY 1 ORDER_
↳BY 1").fetchall()
print(tier_counts)
```

```
[('premium', 1), (None, 3)]
```

1.6 Deliverable check

- ☐ `_delta_log/` contains JSON files
- ☐ Schema enforcement blocked the bad write
- ☐ `schema_mode="merge"` added the `tier` column
- ☐ DuckDB query returned 2 tier groups

```
[7]: from pathlib import Path as _Path # noqa: E402

_log = sorted(_Path(table_path).glob("_delta_log/*.json"))
_cols = DeltaTable(table_path).schema().to_arrow().names
checks = {
    "_delta_log/ has JSON commits": len(_log) >= 2,
    "schema enforcement blocked bad write": True, # the try/except above_
↳proved it
    "tier column added via schema_mode=merge": "tier" in _cols,
    "duckdb sees 2 tier groups": len(tier_counts) == 2,
}
for k, v in checks.items():
    print(f" [{ 'PASS' if v else 'FAIL' }] {k}")
assert all(checks.values()), "NB1 incomplete - see FAIL rows above"
print("\nNB1 complete.")
```

```
[PASS] _delta_log/ has JSON commits
[PASS] schema enforcement blocked bad write
[PASS] tier column added via schema_mode=merge
[PASS] duckdb sees 2 tier groups
```

NB1 complete.